

An empirical study of the naïve REINFORCE algorithm for predictive maintenance of industrial milling machines

Rajesh Siraskar

June 30, 2023

Abstract

In this empirical study, we document the performance of a simple, early reinforcement learning algorithm, REINFORCE, implemented for a predictive maintenance problem – an optimal tool replacement policy for a milling machine. We compare a naïve implementation of REINFORCE against the predictions of industry-grade Stable-Baselines3 (SB3) implementations of three advanced algorithms, namely, Deep Q-Network (DQN), Advantage Actor-Critic (A2C) and Proximal Policy Optimization (PPO). This work is aimed at industrial practitioners not accustomed to the complex hyperparameter tuning often required to get reinforcement learning algorithms to work. Our broad goal was to understand the performance of *untuned* algorithms under various scenarios such (1) simulation-based environment (2) three sets of real tool-wear data (the benchmark IEEE NUAA Ideahouse dataset) (3) increased difficulty level by adding noise levels and a random chance of break-down.

Model performance was measured by how accurately the predictive maintenance agent suggested tool replacement when compared to a deterministic preventive maintenance rule based on the tool-wear threshold. Across variants of the environment, *selected* REINFORCE models demonstrated an incredibly high tool replacement precision of 0.866 against 0.448 for A2C, 0.415 for DQN, and 0.447 for PPO. The recall, F1-score, and F1-beta (0.5) scores were all significantly higher as well. While the selected REINFORCE models demonstrated better

performance for each variant, it was observed that the training was unstable, occasionally producing poor performance models (these were discarded). On the other hand, the SB3 implementations training was more stable, almost always producing models of similar performance significantly low performing models. Our findings indicate that the computationally lightweight REINFORCE performs significantly well for this particular problem. Consequently, for this particular problem, selecting the naïve REINFORCE could be a more suitable and effective policy generating alternative to more advanced complex algorithms.

For reproducibility, model training and testing code, data and the *selected* REINFORCE models have been uploaded to <https://github.com/Link>

Keywords: Predictive maintenance, milling machines, Reinforcement Learning, REINFORCE

Abbreviations

DQN	Deep Q-Network	A2C	Advantage Actor-Critic
PPO	Proximal Policy Optimization		
SS	Single-variable state	MS	Multi-variate state
TP	True positive	TN	True negative
FP	False positive	FN	False negative
RL	Reinforcement Learning	SB3	Stable-Baselines3

1 Introduction

"Plurality should not be posited without necessity" – Of two competing theories, the simpler explanation of an entity is to be preferred

— William of Ockham (1285–1347), *The Occams razor principle*

Milling machines are highly versatile, ubiquitous tools serving a variety of industries. A milling machine removes metal from the work piece by rotating and driving a cutting device into it. Abrasive forces cause tool wear, and optimal tool replacement reduces direct costs and optimizes the machines' downtime. With the 2023 milling machine market valued at USD 68.3 billion ([Future Market Insights, 2023](#)), this is an important goal for the industry. The cutting tool experiences multiple types of wear as it cuts through metal. Tool wear depends in several factors such as the cutting speed, force applied to the tool, lubrication and materials of the work piece and cutting tool.

Reinforcement learning (RL) is an artificial intelligence technique inspired by nature. Fig. 1 ([Sutton and Barto, 2018](#)) shows the RL learning feedback loop. An actor or "agent" interacts with an environment and learns via "trial-and-error". It acts based on stimuli or feedback received from the environment after performing a certain action. Actions that help in achieving the learning goal receive a reward while actions that do not, are punished. Repeating this loop over thousands of episodes, good actions are "reinforced", thereby building a "policy" that is optimized for that goal. In the case of predictive maintenance for milling machines, the agent is the "planner" with a goal of learning an optimal tool replacement policy. The environment consists of sensors attached to the machine and related information such as job specifications, environment conditions etc.

Introduced in 1992, the REINFORCE algorithm ([Williams, 1992](#)) is considered as a basic reinforcement learning algorithm. It is a policy-based, on-policy algorithm, capable of handling both discrete and continuous observation and action domains.

In practice the REINFORCE algorithm is considered as a "weak" learner and superseded by several algorithms developed since. Most notably the Q-Learning and its deep-neural network version, the DQN ([Mnih et al., 2013](#)),

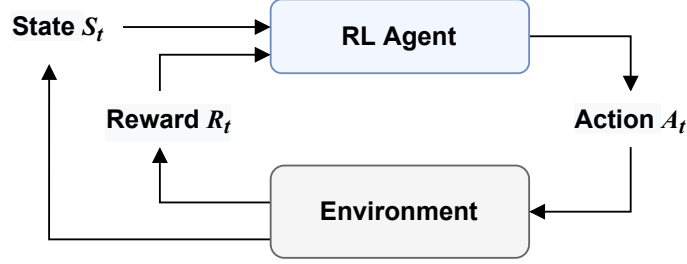


Figure 1: Reinforcement Learning

followed by Actor-Critic (Mnih et al., 2016) and one of the most robust modern-day algorithms, the PPO (Schulman et al., 2017).

1.1 Stable-Baselines3

Stable-Baselines3 Raffin et al. (2021), is Open Source and very popular among the RL community¹. Stable-Baselines3 was initially based on the Open AI baselines (Dhariwal et al., 2017) and is completely rewritten using PyTorch.

for , implementation s the advanced algorithms – DQN, A2C and PPO and compare its performance to a naïve custom implementation of the REINFORCE algorithm. We use Stable-Baselines3 (SB3), the highly popular and reliable implementations of DQN, A2C and PPO. As of 27-Jun-2023, the REINFORCE was not implemented by Stable-Baselines and we therefore custom implemented a basic version.

1.2 REINFORCE

sB-3 dOES NOT IMPLEMENT THE REINFORCE — <https://stable-baselines3.readthedocs.io/en/stable/faq.html#does-stable-baselines3-implement-reinforce>

Model Free -> Policy Gradient/Actor-Critic -> REINFORCE) REINFORCE (Monte-Carlo policy gradient)

IEEE NUAA Ideahouse dataset has been used in this paper for the RUL estimation of the milling cutter. Yingguang et al. (2021)

Tool wear modeling is the first step to assist in predicting

¹As of 27-Jun-2023, it had 6k+ stars and 424 closed pull requests

Literature search conducted on the ScopusTM and Web Of ScienceTM did not return any articles for the application of reinforcement learning for predictive maintenance of milling machines. Search strings we tried – “reinforcement learning AND tool wear AND maintenance”, RL + milling + policy, RL + milling + maintenance, RL + tool wear + policy, RL + tool wear + maintenance"

Running the search “reinforcement learning AND milling AND tool wear" using the ScopusTM and Web Of ScienceTM services

[Dai et al. \(2021\)](#) is the only article we found that tackles the

Machine learning methods have been applied for example – [Oshida et al. \(2023\)](#) proposes real-time tool wear detection during the milling process. They use a stacked LSTM encoder-decoder model for anomaly detection.

No results for "reinforcement learning" AND "milling machine" AND "tool wear" - on scopus or was as of 23-jun-2023 "reinforcement learning" AND "milling machine" - 1 not relevant " *Conference Paper* • Open access "Online Learning of Stability Lobe Diagrams in Milling" Friedrich, J. Torzewski, J. Verl, A.

"reinforcement learning" AND "tool wear" - 10 results

2 Literature Review

3 Method

The methodology explains in detail what the researcher did to undertake the research. Various aspects of the research have to be outlined: The overall structure and operation of the experiment or observational experience. The groups studied in the research including the size of each group and any features of the subjects which may be relevant to the topic being researched. The variables that were changed between groups and the variables measured as a result of the changes. The conditions under which the research was undertaken and any factors or variations in conditions which may have an impact on the results. The methods of data analysis used in order to analyse and

*collate the results. ***Any limitations of the data collected.* We normalize the tool wear and other state features, $x \in [0, 1] \subset \mathbb{R}$. This allows for adding white noise of similar magnitudes across experiments of different data-sets

3.1 RL environment description

3.2 Data description

3.3 Procedure

- training
- selecting the model
- conducting the experiments

3.4 Evaluation method

3.5 Tools

2 different laptops

- why classification metrics
- why F1beta

3.6 Method - training and testing

3.7 Precision or Recall?

- Precision => low FP => False replacement-action reduced. *** Unnecessary tool replacement reduced. Tool life maximized, down time minimized, production disruption minimized
- Recall => low FN => False declaration of normal operation reduced. Reduce missed replacements. Tool replacements increased. *** Product quality not compromised.

3.8 Hyper-parameters for Precision or Recall control

- $R1 = +1$
- $R2 = -1$
- $R3 = -100 \Rightarrow$ higher neg. Improve recall. Lower neg. Improve precision
- **LOOK AHEAD PARAM:**
- Training: SB3 - 10 k eps. 3 times. Average their outputs
- Testing:
 - Avg. over 5 rounds.
 - Each round - avg over 40 test cases x 10 test rounds
 - Total: $40 \times 10 \times 5 = 2000$ cases
 - Avgs over: 10 rounds (of 40 cases each) X 5 rounds of **re-trained** SB3 agents = 50 rounds

3.9 Inference

- Training: SB3 is also unstable - show examples of results such as A2C/DQN 0.00
- Training: SB3 is also unstable - SHOW SB3 tensorboard plots
- Training: SB3 is also unstable - EXCEL plots of results over the 10 rounds

4 Network architecture and basic hyper-parameters

source of ppo implementation details <https://iclr-blog-track.github.io/2022/03/25/ppo-implementation-details/>

source of SB3 network mp : <https://github.com/openai/baselines/blob/ea25b9e8b234e6ee1bca43083f8f3cf974143998/baselines/common/models.py#L75-L103>

	A2C	DQN	PPO	REINFORCE
Network architecture	input dim x [64 Tanh x 64 Tanh] x output dim	input dim x [64 Tanh x 64 Tanh] x output dim	input dim x [64 Tanh x 64 Tanh] x output dim	input dim x [64 ReLU] x output dim
Layers	2	2	2	1
Units	64 x 64	64 x 64	64 x 64	64
Activation	Tanh, Tanh	Tanh, Tanh	Tanh, Tanh	ReLU
Optimizer	RMSprop	Adam	Adam	Adam
Learning rate	0.0007	0.0001	0.0003	0.01
Gamma	0.99	0.99	0.99	0.99

Table 1: Comparing the network architecture and basic hyper-parameters across algorithms

4.1 PPO hyperparms

implementation guide source » <https://iclr-blog-track.github.io/2022/03/25/ppo-implementation-details/>

By default, PPO uses a simple MLP network consisting of two layers of 64 neurons and Hyperbolic Tangent as the activation function. Then PPO builds a policy head and value head that share the outputs of the MLP network. Below is a pseudocode:


4.2 dqn hyperparms

default hyperparms : https://stable-baselines3.readthedocs.io/en/master/_modules/stable_baselines3/common/policies.html


overridden in indiv policies for example SB3 DQN hyperparms for example were taken from ; Paper: <https://arxiv.org/abs/1312.5602>, <https://www.nature.com/articles/> Default hyperparameters are taken from the Nature paper, except for the optimizer and learning rate that were taken from Stable Baselines defaults

5 Empirical results


5.1 Detailed metrics




(a) Precision



(b) Recall




(c) F1-score




(d) F1-beta (0.5)


Figure 2: Overall performance – Metrics across 10 rounds of model training




(a) Precision



(b) Recall




(c) F1-score




(d) F1-beta (0.5)


Figure 3: Simulated environment – Metrics across 10 rounds of model training



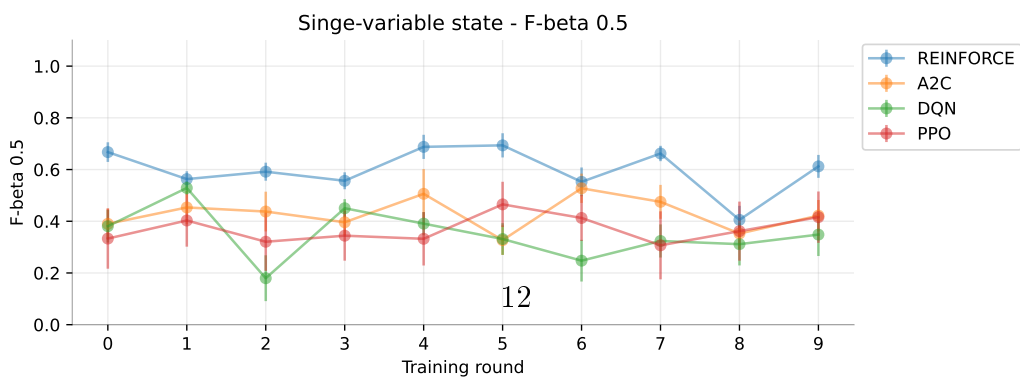
(a) Precision



(b) Recall




(c) F1-score




(d) F1-beta (0.5)


Figure 4: Singe-variable state environment – Metrics across 10 rounds of model training




(a) Precision



(b) Recall



(c) F1-score



(d) F1-beta (0.5)

Figure 5: Multi-variate state environment – Metrics across 10 rounds of model training




Figure 6: Training time. Across 10 rounds and all environment variants.

Environment	REINFORCE				A2C				DQN				PPO			
	Prec.	Recall	F1	F0.5	Prec.	Recall	F1	F0.5	Prec.	Recall	F1	F0.5	Prec.	Recall	F1	F0.5
Simulated - No noise	0.999	0.645	0.782	0.898	0.335	0.359	0.344	0.338	0.348	0.597	0.410	0.352	0.392	0.211	0.252	0.30
Simulated - Low noise	0.943	0.954	0.948	0.945	0.409	0.318	0.349	0.379	0.273	0.064	0.076	0.108	0.359	0.173	0.205	0.25
Simulated - High noise	0.889	0.974	0.929	0.904	0.471	0.439	0.443	0.455	0.423	0.408	0.295	0.284	0.402	0.205	0.248	0.30
PHM C01 SS - No noise	0.886	0.978	0.928	0.902	0.294	0.337	0.305	0.296	0.350	0.405	0.291	0.269	0.517	0.494	0.471	0.47
PHM C01 SS - Low noise	0.916	0.893	0.903	0.911	0.526	0.645	0.568	0.540	0.321	0.591	0.404	0.343	0.490	0.415	0.443	0.46
PHM C01 SS - High noise	0.757	0.926	0.831	0.784	0.499	0.632	0.542	0.513	0.399	0.402	0.308	0.292	0.403	0.223	0.270	0.32
PHM C04 SS - No noise	0.865	0.959	0.908	0.881	0.515	0.676	0.575	0.535	0.365	0.497	0.383	0.348	0.431	0.239	0.265	0.31
PHM C04 SS - Low noise	0.722	0.980	0.831	0.762	0.399	0.393	0.391	0.393	0.409	0.589	0.410	0.361	0.438	0.299	0.334	0.37
PHM C04 SS - High noise	0.770	0.809	0.787	0.776	0.375	0.456	0.397	0.381	0.408	0.411	0.296	0.282	0.491	0.324	0.362	0.40
PHM C06 SS - No noise	0.996	0.609	0.751	0.879	0.463	0.454	0.455	0.459	0.538	0.780	0.585	0.523	0.402	0.410	0.374	0.37
PHM C06 SS - Low noise	0.968	0.854	0.905	0.941	0.508	0.615	0.548	0.522	0.395	0.593	0.411	0.362	0.454	0.342	0.367	0.40
PHM C06 SS - High noise	0.699	0.912	0.790	0.732	0.480	0.512	0.466	0.467	0.581	0.499	0.417	0.433	0.424	0.199	0.252	0.31
PHM C01 MS - No noise	0.824	0.895	0.856	0.836	0.444	0.284	0.315	0.358	0.313	0.215	0.165	0.175	0.513	0.347	0.395	0.44
PHM C04 MS - No noise	0.752	0.678	0.709	0.733	0.506	0.326	0.368	0.425	0.588	0.642	0.492	0.486	0.472	0.455	0.444	0.45
PHM C06 MS - No noise	1.000	0.643	0.779	0.896	0.499	0.731	0.575	0.523	0.520	0.239	0.209	0.256	0.509	0.260	0.330	0.40

Table 2: Model performance comparison all variants of the environments.

5.2 Overall summary performance

	Precision		Recall		F1-score		F1-beta score	
	Mean	SD	Mean	SD	Mean	SD	Mean	SD
A2C	0.448	0.074	0.478	0.084	0.443	0.071	0.439	0.069
DQN	0.415	0.196	0.462	0.033	0.343	0.038	0.325	0.063
PPO	0.447	0.147	0.306	0.090	0.334	0.093	0.375	0.107
REINFORCE	0.866	0.042	0.847	0.054	0.842	0.043	0.852	0.042

Table 3: Model performance summary - averaged over all environment.

5.3 Simulated environment

	Precision		Recall		F1-score		F1-beta score	
	Mean	SD	Mean	SD	Mean	SD	Mean	SD
A2C	0.405	0.079	0.372	0.086	0.379	0.076	0.391	0.076
DQN	0.348	0.217	0.356	0.033	0.260	0.041	0.248	0.068
PPO	0.385	0.175	0.196	0.064	0.235	0.080	0.289	0.110
REINFORCE	0.944	0.029	0.858	0.041	0.886	0.032	0.916	0.030

Table 4: Model performance summary - averaged over simulated environments.

5.4 Real data – simple single-variable environment

	Precision		Recall		F1-score		F1-beta score	
	Mean	SD	Mean	SD	Mean	SD	Mean	SD
A2C	0.451	0.064	0.524	0.085	0.472	0.067	0.456	0.063
DQN	0.418	0.172	0.530	0.032	0.389	0.034	0.357	0.055
PPO	0.450	0.146	0.327	0.095	0.349	0.095	0.384	0.106
REINFORCE	0.842	0.043	0.880	0.053	0.848	0.043	0.841	0.042

Table 5: Model performance summary - averaged over PHM-2010 environments with simple single-variable environment.

5.5 Real data – complex multi-variate environment

	Precision		Recall		F1-score		F1-beta score	
	Mean	SD	Mean	SD	Mean	SD	Mean	SD
A2C	0.483	0.101	0.447	0.081	0.419	0.075	0.435	0.079
DQN	0.474	0.248	0.365	0.038	0.289	0.049	0.306	0.082
PPO	0.498	0.121	0.354	0.103	0.390	0.101	0.436	0.104
REINFORCE	0.859	0.053	0.739	0.069	0.781	0.055	0.822	0.052

Table 6: Model performance summary - averaged over PHM-2010 environments with complex multi-variate environment.

5.6 Training times

6 The REINFORCE algorithm

Three key features of any RL algorithm:

Environment	REINFORCE	A2C	DQN	PPO
Simulated - No noise	214.23	41.19	4.03	41.13
Simulated - Low noise	199.89	41.52	3.55	40.66
Simulated - High noise	134.16	17.88	1.53	20.90
PHM C01 SS - No noise	330.54	18.85	2.08	32.65
PHM C01 SS - Low noise	426.79	30.66	3.69	38.59
PHM C01 SS - High noise	333.13	17.58	1.80	19.16
PHM C04 SS - No noise	299.31	19.56	1.86	19.64
PHM C04 SS - Low noise	264.90	18.27	2.00	19.69
PHM C04 SS - High noise	256.44	17.65	1.58	19.11
PHM C06 SS - No noise	339.65	17.64	2.26	19.50
PHM C06 SS - Low noise	266.98	19.33	1.84	19.19
PHM C06 SS - High noise	308.20	34.21	4.18	30.94
PHM C01 MS - No noise	655.21	38.55	4.96	42.21
PHM C04 MS - No noise	615.58	33.85	7.36	43.49
PHM C06 MS - No noise	625.37	39.30	5.85	41.68
Overall average (s)	351.36	27.07	3.24	29.90

Table 7: Training time for each model, across different environments. Averaged over 3 runs. Time is in seconds (s).

1. Policy: π_θ = Probabilities of all actions, given a state. Parameterized by θ
2. Objective function:

$$\max_{\theta} J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)] \quad (1)$$

3. Method: Way to update the parameters = Policy Gradient

6.1 Policy gradient numerical computation

1. Plain vanilla:

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T R_t(\tau) \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \right] \quad (2)$$

2. With Monte Carlo sampling and approximation: $\nabla_{\theta} J(\pi_{\theta}) \approx \left[\sum_{t=0}^T R_t(\tau) \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \right]$
3. With baseline: $\nabla_{\theta} J(\theta) \approx \left[\sum_{t=0}^T (R_t(\tau) - b(s_t)) \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) \right]$
4. Where, baseline does not change per time-step, it is for the entire trajectory
5. One baseline option: V^{π} - leads to Actor-Critic algorithm
6. Simpler option: Average returns over trajectory: $b = \frac{1}{T} \sum_{t=0}^T R_t(\tau)$

Algorithm

7 Discussion

7.1 Precision or Recall?

- Precision \Rightarrow low FP \Rightarrow False replacement-action reduced. *** Unnecessary tool replacement reduced. Tool life maximized, down time minimized, production disruption minimized
- Recall \Rightarrow low FN \Rightarrow False declaration of normal operation reduced. Reduce missed replacements. Tool replacements increased. *** Product quality not compromised.

limitations of the selecting model for evaluation

[Raffin et al. \(2021\)](#) – quote from paper - intro section – "A major challenge is that small implementation details can have a substantial effect on performance often greater than the difference between algorithms (Engstrom et al., 2020). It is particularly important that implementations used as experimental baselines are reliable; otherwise, novel algorithms compared to weak baselines lead to inflated estimates of performance improvements"

DPO paper - Rafael Rafailov May-2023 — Direct Preference Optimization: Your Language Model is Secretly a Reward Model

5.2 Instability of Actor-Critic Algorithms We can also use our framework to diagnose instabilities with standard actor-critic algorithms used for the RLHF, such as PPO. We follow the RLHF pipeline and focus on the RL fine-tuning step outlined in Section 3. We can draw connections to the control as inference framework [18] for the constrained RL problem outlined in 3. We assume a parameterized model and minimize This is the same objective optimized in prior works [45, 35, 1, 23] using the DPO-equivalent reward for the reward class of In this setting, we can interpret the normalization term in as the soft value function of the reference policy . While this term does not affect the optimal solution, without it, the policy gradient of the objective could have high variance, making learning unstable. We can accommodate for the normalization term using a learned value function, but that can also be difficult to optimize. Alternatively, prior works have normalized rewards using a human completion baseline, essentially a single sample Monte-Carlo estimate of the normalizing term. In contrast the DPO reparameterization yields a reward function that does not require any baselines.

8 Conclusion

References

- Zixiang Dai, Mingyan Jiang, Xiaogang Li, Dongfeng Yuan, and Xiaotian Zhou. Reinforcement lion swarm optimization algorithm for tool wear prediction. In *2021 Global Reliability and Prognostics and Health Management (PHM-Nanjing)*, pages 1–7. IEEE, 2021.
- Prafulla Dhariwal, Christopher Hesse, Oleg Klimov, Alex Nichol, Matthias Plappert, Alec Radford, John Schulman, Szymon Sidor, Yuhuai Wu, and Peter Zhokhov. Openai baselines. <https://github.com/openai/baselines>, 2017.
- Future Market Insights. Milling machine market outlook (2023 to

- 2033), 1 2023. URL <https://www.futuremarketinsights.com/reports/milling-machine-market>. Accessed: 2023-06-23.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin Riedmiller. Playing atari with deep reinforcement learning. *arXiv preprint arXiv:1312.5602*, 2013.
- Volodymyr Mnih, Adria Puigdomenech Badia, Mehdi Mirza, Alex Graves, Timothy Lillicrap, Tim Harley, David Silver, and Koray Kavukcuoglu. Asynchronous methods for deep reinforcement learning. In *International conference on machine learning*, pages 1928–1937. PMLR, 2016.
- Taisuke Oshida, Tomohiro Murakoshi, Libo Zhou, Hirotaka Ojima, Kazuki Kaneko, Teppei Onuki, and Jun Shimizu. Development and implementation of real-time anomaly detection on tool wear based on stacked lstm encoder-decoder model. *The International Journal of Advanced Manufacturing Technology*, pages 1–16, 2023.
- Antonin Raffin, Ashley Hill, Adam Gleave, Anssi Kanervisto, Maximilian Ernestus, and Noah Dormann. Stable-baselines3: Reliable reinforcement learning implementations. *J. Mach. Learn. Res.*, 22(1), jan 2021. ISSN 1532-4435.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Richard Sutton and Andrew Barto. *Reinforcement Learning: An Introduction*. The MIT Press, Cambridge, England, 2nd. edition edition, 2018.
- Ronald J Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Reinforcement learning*, pages 5–32, 1992.
- LI Yingguang, LIU Changqing, LI Dehua, HUA Jiaqi, and WAN Peng. Tool wear dataset of nuaa ideahouse, 2021. URL <https://dx.doi.org/10.21227/3aa1-5e83>.